

This notebook is an exercise in the [Intermediate Machine Learning](#) course. You can reference the tutorial at [this link](#).

In this exercise, you will leverage what you've learned to tune a machine learning model with **cross-validation**.

Setup

The questions below will give you feedback on your work. Run the following cell to set up the feedback system.

```
In [1]: # Set up code checking
import os
if not os.path.exists("../input/train.csv"):
    os.symlink("../input/home-data-for-ml-course/train.csv", "../input/train.csv")
    os.symlink("../input/home-data-for-ml-course/test.csv", "../input/test.csv")
from learntools.core import binder
binder.bind(globals())
from learntools.ml_intermediate.ex5 import *
print("Setup Complete")
```

Setup Complete

You will work with the [Housing Prices Competition for Kaggle Learn Users](#) from the previous exercise.



Run the next code cell without changes to load the training and test data in `X` and `X_test`. For simplicity, we drop categorical variables.

```
In [2]: import pandas as pd
from sklearn.model_selection import train_test_split

# Read the data
train_data = pd.read_csv('../input/train.csv', index_col='Id')
test_data = pd.read_csv('../input/test.csv', index_col='Id')

# Remove rows with missing target, separate target from predictors
train_data.dropna(axis=0, subset=['SalePrice'], inplace=True)
y = train_data.SalePrice
train_data.drop(['SalePrice'], axis=1, inplace=True)

# Select numeric columns only
numeric_cols = [cname for cname in train_data.columns if train_data[cname].dtype in ['i
```

```
X = train_data[numeric_cols].copy()
X_test = test_data[numeric_cols].copy()
```

Use the next code cell to print the first several rows of the data.

```
In [3]: X.head()
```

```
Out[3]:
```

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	MasVnrArea
Id								
1	60	65.0	8450	7	5	2003	2003	196.0
2	20	80.0	9600	6	8	1976	1976	0.0
3	60	68.0	11250	7	5	2001	2002	162.0
4	70	60.0	9550	7	5	1915	1970	0.0
5	60	84.0	14260	8	5	2000	2000	350.0

5 rows × 36 columns

So far, you've learned how to build pipelines with scikit-learn. For instance, the pipeline below will use `SimpleImputer()` to replace missing values in the data, before using `RandomForestRegressor()` to train a random forest model to make predictions. We set the number of trees in the random forest model with the `n_estimators` parameter, and setting `random_state` ensures reproducibility.

```
In [4]: from sklearn.ensemble import RandomForestRegressor
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer

my_pipeline = Pipeline(steps=[
    ('preprocessor', SimpleImputer()),
    ('model', RandomForestRegressor(n_estimators=50, random_state=0))
])
```

You have also learned how to use pipelines in cross-validation. The code below uses the `cross_val_score()` function to obtain the mean absolute error (MAE), averaged across five different folds. Recall we set the number of folds with the `cv` parameter.

```
In [5]: from sklearn.model_selection import cross_val_score

# Multiply by -1 since sklearn calculates *negative* MAE
scores = -1 * cross_val_score(my_pipeline, X, y,
                              cv=5,
                              scoring='neg_mean_absolute_error')

print("Average MAE score:", scores.mean())
```

Average MAE score: 18276.410356164386

Step 1: Write a useful function

In this exercise, you'll use cross-validation to select parameters for a machine learning model.

Begin by writing a function `get_score()` that reports the average (over three cross-validation folds) MAE of a machine learning pipeline that uses:

- the data in `X` and `y` to create folds,
- `SimpleImputer()` (with all parameters left as default) to replace missing values, and
- `RandomForestRegressor()` (with `random_state=0`) to fit a random forest model.

The `n_estimators` parameter supplied to `get_score()` is used when setting the number of trees in the random forest model.

```
In [6]: def get_score(n_estimators):
        """Return the average MAE over 3 CV folds of random forest model.

        Keyword argument:
        n_estimators -- the number of trees in the forest
        """
        my_pipeline = Pipeline(steps = [('preprocessor', SimpleImputer()), ('model', RandomForestRegressor(n_estimators=n_estimators))]
        # Replace this body with your own code
        scores = -1 * cross_val_score(my_pipeline, X, y, cv = 3, scoring = 'neg_mean_absolute_error')
        return scores.mean()

        # Check your answer
        step_1.check()
```

Correct

```
In [7]: # Lines below will give you a hint or solution code
        #step_1.hint()
        #step_1.solution()
```

Step 2: Test different parameter values

Now, you will use the function that you defined in Step 1 to evaluate the model performance corresponding to eight different values for the number of trees in the random forest: 50, 100, 150, ..., 300, 350, 400.

Store your results in a Python dictionary `results`, where `results[i]` is the average MAE returned by `get_score(i)`.

```
In [8]: results = {}
        for i in range(1,9): results[i*50] = get_score(50*i)
```

```
# Check your answer
step_2.check()
```

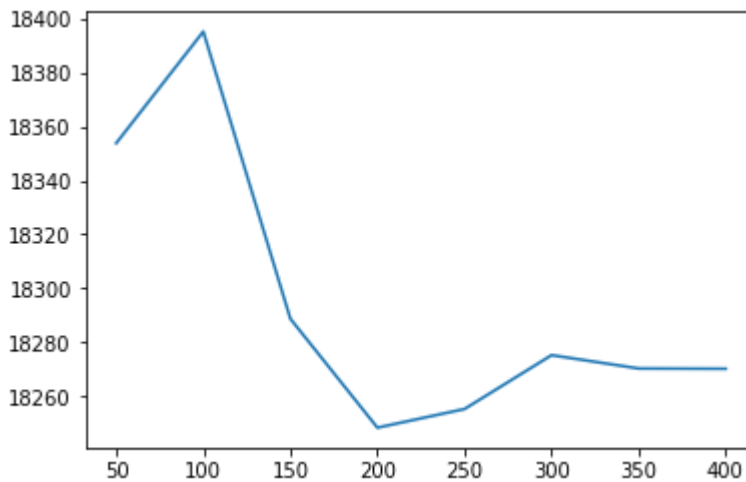
Correct

```
In [9]: # Lines below will give you a hint or solution code
        #step_2.hint()
        #step_2.solution()
```

Use the next cell to visualize your results from Step 2. Run the code without changes.

```
In [10]: import matplotlib.pyplot as plt
         %matplotlib inline

         plt.plot(list(results.keys()), list(results.values()))
         plt.show()
```



Step 3: Find the best parameter value

Given the results, which value for `n_estimators` seems best for the random forest model? Use your answer to set the value of `n_estimators_best`.

```
In [11]: n_estimators_best = min(results, key=results.get) # get the minimum from the dict

        # Check your answer
        step_3.check()
```

Correct

```
In [12]: # Lines below will give you a hint or solution code
         #step_3.hint()
         #step_3.solution()
```

In this exercise, you have explored one method for choosing appropriate parameters in a machine learning model.

If you'd like to learn more about [hyperparameter optimization](#), you're encouraged to start with **grid search**, which is a straightforward method for determining the best *combination* of parameters for a machine learning model. Thankfully, scikit-learn also contains a built-in function `GridSearchCV()` that can make your grid search code very efficient!

Keep going

Continue to learn about [gradient boosting](#), a powerful technique that achieves state-of-the-art results on a variety of datasets.

Have questions or comments? Visit the [course discussion forum](#) to chat with other learners.